

Standard-Cell Characterization Methods for sg13g2_stdcell_hv



fo4.py (reference) · CharLib 2.1.0 · lctime 0.0.26 — the measurement physics behind the Liberty file, and the cells no characterizer reaches

Koen Van Caekenberghe, Ph.D.

ChipDesign B.V.

info@chipdesign.be

2026-08-18 (rev. 2: special-cell classes, bus-holder capacitance, drive limits)

1 Scope and headline result

Three tools touched the sg13g2_stdcell_hv Liberty file, in three different roles:

- **work/fo4.py** — the *reference anchor*: a 96-line ngspice deck that measures the FO4 delay and the input capacitance of the thick-oxide inv_1 against the thin-oxide original. Its two ratios (2.66× delay, 2.20× capacitance) set the slew and load grids of the shipped library, and its rail-biased AC capacitance (5.87 fF) is the reference every other capacitance number in this report is judged against.
- **CharLib 2.1.0** (commit 6859faf, driven through the infinitymdm PySpice 1.6 fork) — the *production characterizer* that generated the shipped NLDM tables, configured with the charge_integration input- capacitance procedure and, for the sequential cells, the project's own seq_delay_procedure.py registered through CharLib's procedure registry.
- **lctime 0.0.26** (LibreCell) — the *independent cross-check*: work/lctime_compare.py re-characterized eight combinational cells on identical grids and models and aligned 3 132 table points against the shipped library.

Every method claim below was verified in the installed source of the three tools (file and line references throughout; a provenance appendix closes the report). The comparison in one table, before the derivations:

quantity	fo4.py (reference)	CharLib 2.1.0 as configured	lctime 0.0.26
delay / slew load	four real inverter gates	lumped capacitor	lumped capacitor
stimulus ramp	shaped by a driver stage	PWL, stretched by 1/0.6 (Liberty-correct)	PWL equal to the raw slew number (convention error)

extraction

quantity	fo4.py (reference)	CharLib 2.1.0 as configured	lctime 0.0.26
	ngspice .meas trig/ targ	ngspice .meas trig/ targ	numpy interpolation of waveforms
input capacitance	AC at 1 MHz, rail- biased, mean of both rails	charge integration, worst edge	constant-current slope 10 μ A, 20–80 % window
inv_1 C _{in} result	5.87 fF	6.44 fF (+9.7 %)	8.92 fF (+52 %)
internal power	not measured	not implemented in CharLib 2.1.0	rectangle-rule energy integration (load energy not subtracted)
leakage	not measured	DC operating point per input state, all 2 ^N states	unimplemented stub
setup / hold	—	pass/fail bisection, 1.5× C2Q degradation (local procedure)	Brent root-finding, 10 ps absolute pushout
clock→Q	—	local procedure (upstream is a stub)	implemented; constraint search at zero output load
simulator coupling	ngspice -b batch	libngspice in-process (shared)	ngspice subprocess, ASCII wrdata

Answer to the standing question — can lctime use CharLib’s charge-integration method? Not as shipped: lctime 0.0.26 contains no code path that integrates a current anywhere (no trapz, no cumsum, no ngspice integ in any generated .control block). Its input capacitance is a constant-current secant slope, which is where its +52 % error comes from. However, lctime’s simulator coupling *already* returns branch currents with the time vector — it integrates supply and gate current for its internal-power tables — so a charge-integration capacitance is a localized, three-part patch to one function. Section 8 gives the exact changes.

2 What the Liberty NLDM asks for

All three tools ultimately serve the same data model, so the definitions come first. The library header fixes eight trip points; the shipped file uses the industry-standard set

$$V_{th,in} = V_{th,out} = 0.5 V_{DD}, \quad V_{slew,lo} = 0.2 V_{DD}, \quad V_{slew,hi} = 0.8 V_{DD}$$

and defines, per timing arc, the propagation delay and the output transition time as threshold-crossing intervals of the simulated waveforms:

$$t_{pd} = t(v_{out} = 0.5 V_{DD}) - t(v_{in} = 0.5 V_{DD})$$

$$t_{tran} = t(v_{out} = 0.8 V_{DD}) - t(v_{out} = 0.2 V_{DD})$$

(rising conventions shown; falling edges mirror the thresholds). The NLDM `cell_rise / rise_transition` groups tabulate these two quantities as functions of two independent variables,

$$t_{pd} = f(s_{in}, C_L), \quad t_{tran} = g(s_{in}, C_L),$$

where s_{in} is the *input* transition time measured between the same 20 %/80 % slew thresholds, and C_L the lumped load. The shipped library uses 7×7 grids per arc; the grid values are the thin-oxide library's grids scaled by the fo4.py ratios so the tables span the same electrical territory. Everything a characterizer does for a combinational arc reduces to filling f and g point by point from transient simulations, and every methodological difference between the three tools is a difference in *stimulus construction*, *load modelling*, or *waveform extraction* for those same two equations.

A physical scale for the delays themselves comes from logical effort: with $\tau = R_{inv}C_{in}$ the technology time constant, a stage driving an electrical effort $h = C_{out}/C_{in}$ has normalized delay

$$d = g h + p,$$

and the fanout-of-4 inverter delay ($g = 1$, $h = 4$) is the canonical technology speed metric fo4.py measures directly: $t_{FO4} = \tau (4 + p)$. It is a *ratio-preserving* metric — sizing every cell by the same factors leaves g and p nearly unchanged — which is exactly why one FO4 number per library (142.4 ps HV vs 53.6 ps LV) legitimately rescales an entire grid set.

3 Input capacitance: three answers to one question

3.1 The physics being sampled

The input pin of a CMOS gate is not a capacitor; it is a nonlinear, *transcapacitive* charge reservoir. The gate charge is a function of the gate voltage **and** of every other terminal voltage, so the differential capacitance seen while the input traverses its switching trajectory is

$$C_{in}(v) = \frac{dQ_G}{dv_{in}} = C_{gs}(v) + C_{gb}(v) + C_{gd}(v) \left(1 - \frac{dv_{out}}{dv_{in}}\right).$$

The last term is the Miller effect: while the cell switches, the output moves *against* the input with the incremental gain $A_v = dv_{out}/dv_{in} < 0$, so C_{gd} contributes $C_{gd}(1 + |A_v|)$ — a sharp peak centred on the switching threshold V_M , where $|A_v|$ is largest. Away from V_M , with the output pinned at a rail, $A_v \approx 0$ and C_{in} relaxes to its small end-point values.

Every “input capacitance” is therefore a *functional* of $C_{in}(v)$ — a weighted average over some window of the input swing — and the three tools choose three different windows. That single observation explains all three numbers.

3.2 fo4.py: small-signal AC at the rails

fo4.py drives an isolated inverter input with a 1 V AC source at 1 MHz, once biased at $v = 0$ and once at $v = V_{DD}$ (fo4.py:38–55). For a linear one-port $Y(j\omega) = G + j\omega C$, the measured current at unit drive gives

$$C_{AC} = \frac{\text{Im}\{I\}}{2\pi f |\hat{V}|} \quad (\text{evaluated at } f = 1 \text{ MHz}).$$

The frequency is chosen so that gate leakage is invisible ($G \ll \omega C$: at 1 MHz, $\omega C \approx 3.7 \times 10^{-8}$ S for 5.9 fF, orders above any gate conductance) yet low enough that non-quasi-static channel effects play no role. The deck comments state the bias rationale explicitly: at either rail the inverter sits in a zero-gain region, so $A_v \approx 0$ and the Miller term vanishes —

$$C_{AC} = \frac{1}{2} [C_{in}(0) + C_{in}(V_{DD})] \quad (\text{endpoint sample, no Miller}).$$

A mid-rail bias would instead sample the peak of $C_{in}(v)$ and Miller-multiply C_{gd} into the answer. Result: **5.87 fF** for sg13g2_hv_inv_1.

3.3 CharLib: charge integration

The configured procedure (charlib/characterizer/procedures/pin_capacitance/charge_integration.py) drives the pin with a PWL *voltage* ramp — VSS up to VDD, wait, back down — and lets ngspice integrate the stimulus-source branch current over each edge window:

```
.meas TRAN q_rise integ i(vstim) from=... to=...
.meas TRAN q_fall integ i(vstim) from=... to=...
```

then applies the defining relation of effective capacitance (charge_integration.py:100–135):

$$C_{int} = \frac{|\Delta Q|}{\Delta V} = \frac{1}{V_{DD}} \int_{\text{edge}} i_{in}(t) dt = \frac{1}{V_{DD}} \int_0^{V_{DD}} C_{in}(v) dv.$$

This is the *full-swing mean* of $C_{in}(v)$ along the real switching trajectory: the Miller charge $Q_M \approx C_{gd}(\Delta v_{in} + \Delta v_{out})$ is included in ΔQ , but it is spread over the full denominator V_{DD} . Non-driven pins are isolated with $10 \text{ G}\Omega \parallel 1 \text{ pF}$; the ramp time defaults to the fastest grid slew and the wait time to 1000× that, so each edge integrates to completion. The reported capacitance is the worst (larger) of the two edges. Result: **6.44 fF** (+9.7 % vs the AC reference — that surplus *is* the Miller and nonlinearity charge, not an error; it is the charge a driving cell genuinely must deliver across a full transition).

Worth recording: CharLib's *upstream default* is a different procedure, `ac_sweep`, which fits capacitance as the slope of the admittance *magnitude* versus frequency, $C = d|Y|/df$. Since $|Y| = 2\pi fC$ for a capacitor, the correct estimator is $C = \frac{1}{2\pi} d|Y|/df$; the missing $1/2\pi$ makes the default overestimate by $\approx 6.28\times$, and the observed 6.8–7.5× (`inv_1`: 43.8 fF) is exactly 2π compounded with the Miller contribution of its mid-transition drive. This is why the project configuration selects `charge_integration`.

3.4 `lctime`: constant-current secant slope

`lctime` does something else entirely (`lctime/characterization/input_capacitance.py`, docstring: “*Measurement of the input capacitance by driving the input pin with a constant current*”). A fixed source $I = 10\ \mu\text{A}$ (hard-coded, `util.py`:162) charges the pin; the pin *voltage* alone is recorded; and the capacitance is the secant slope between the two slew trip points (`input_capacitance.py`:273–297):

$$C_{\text{slope}} = \frac{I}{\Delta V/\Delta t} = \frac{I[t(0.8V_{DD}) - t(0.2V_{DD})]}{0.6V_{DD}}.$$

From first principles this is again a windowed mean of the same $C_{in}(v)$: with $C_{in}(v) dv/dt = I$,

$$\Delta t = \frac{1}{I} \int_{0.2V_{DD}}^{0.8V_{DD}} C_{in}(v) dv \Rightarrow C_{\text{slope}} = \frac{1}{0.6V_{DD}} \int_{0.2V_{DD}}^{0.8V_{DD}} C_{in}(v) dv.$$

The window $[0.2, 0.8] V_{DD}$ *contains* the Miller peak at V_M and *excludes* the low-capacitance tails near the rails; the Miller charge is divided by $0.6 V_{DD}$ instead of V_{DD} . Both effects push the same direction, which orders the three estimators from first principles alone — the Miller-free endpoint sample below the full-swing mean below the peak-window mean:

$$C_{AC} < C_{int} < C_{\text{slope}}$$

— and the measurements land in exactly that order: $5.87 < 6.44 < 8.92$ fF, the slope method **+52 %** above the reference. `lctime` averages the result over all 2^n static states of the other inputs and over both edge directions (reduction per `calc-mode`: mean for typical), and carries a genuine bug while doing so: the returned dictionary swaps the labels, assigning the falling measurement to `rise_capacitance` and vice versa (`input_capacitance.py`:318–321). The default capacitance attribute, being the mean of both, hides the swap.

4 Delay and output slew

4.1 Stimulus: the slew convention is where the tools diverge

Liberty defines s_{in} between the 20 %/80 % thresholds. A characterizer that drives the pin with a single linear ramp of total duration T ($0 \rightarrow 100\%$) realizes a Liberty slew of

$$s = (0.8 - 0.2)T = 0.6T \quad \Longleftrightarrow \quad T = \frac{s}{0.6}.$$

CharLib gets this right. Its `slew_pwl` helper (`utils.py:49–65`, explicitly citing the Liberty User Guide) stretches the requested slew to the full ramp, $T = s/(V_{hi} - V_{lo}) = s/0.6$, before building the PWL source.

lctime does not. Its `StepWave` for combinational stimulus is built with `rise_threshold=0`, `fall_threshold=1` (`timing_combinatorial.py:186–191`), i.e. the table's slew number is used directly as the 0→100 % ramp time. The circuit therefore sees an edge whose *Liberty* slew is only 0.6 s — every `lctime` table point at index s is actually a measurement at 0.6 s . Since $\partial t_{pd}/\partial s_{in} > 0$, `lctime` under-reads delays, negligibly at fast edges and severely at slow ones. The cross-check quantifies it: over the STA-relevant region the two tools agree to a median 2.9 % on delays and 0.0 % on output transitions (3 132 points, 8 cells), but at the 3.36 ns slew point of `inv_1` (0.396 pF) a direct ngspice measurement gives 2.3605 ns where CharLib's table says 2.3571 ns (**+0.15 %**) and `lctime` says 1.9158 ns (**−19 %**); re-interpolating the CharLib table at 0.6 s reproduces `lctime` to 0.25 %, nailing the mechanism. Ironically `lctime` *reads* all eight trip points from the Liberty template header (`util.py:377–398`) and measures with them correctly — only the stimulus construction ignores them.

`fo4.py` sidesteps the question: its device under test is driven by a real inverter stage (`Xdrv`), so the edge shape is whatever the technology produces, and the measured 50 % crossings use ngspice `.meas` with `rise=2/fall=2` — the *second* edge, past the initial-condition transient, so the measurement is taken in periodic steady state.

4.2 Load modelling

Both characterizers drive a lumped capacitor: C_L to ground on the output (`CharLib delay.py:83`; `lctime` deck `Cload` elements). `fo4.py` loads the DUT with **four actual inverter inputs** — nonlinear $C_{in}(v)$ loads that also kick Miller charge back into the driving node. A lumped capacitor equal to $4 C_{in}$ is the NLDM abstraction of that load; the FO4 deck is the ground truth the abstraction is checked against. This is precisely why the library's load grid is anchored to a *measured* C_{in} ratio rather than a nominal one.

Side-input handling differs in a way that shows up in multi-input cells: CharLib simulates every non-masking static state of the other inputs and records the **worst case**, `criterion = max` (`delay.py:16`, 188–190); `lctime` enumerates the unateness-consistent states and reduces with the `calc-mode` function — `np.mean` under the default `typical` mode (`timing_combinatorial.py:77–81`). Two tables built from identical simulations can therefore legitimately differ on any cell where the arc delay depends on the side-input state (stack position effects).

4.3 Extraction

`fo4.py` and CharLib let ngspice do the measurement — `.meas tran ... trig v(in) val=... targ v(out) val=...` — which interpolates crossings inside the simulator on the adaptive time grid. `lctime` post-processes: it normalizes the stored waveforms to $[0, 1]$, mirrors falling edges, and linearly interpolates the bracketing samples of each threshold crossing in numpy (`util.py:251–374`). Both are first-order interpolations of the same transient;

the difference is not accuracy but *plumbing* — and the plumbing matters, as §7 shows: PySpice’s subprocess backend silently drops `.meas` results, which is why CharLib must run in shared-library mode, while lctime never uses `.meas` at all and is immune.

5 Energy, power, leakage

5.1 The energy bookkeeping a Liberty file expects

Per rising output transition the supply delivers charge $Q = \int i_{VDD} dt$, hence energy

$$E_{supply} = V_{DD} \int i_{VDD}(t) dt = V_{DD} \Delta Q.$$

Of this, $C_L V_{DD}^2$ is associated with the external load (half stored on C_L , half dissipated in the pull-up network), and the remainder is *internal*: parasitic-node charging plus the short-circuit (crowbar) charge that flows while both networks conduct near V_M . Liberty’s `internal_power` tables are meant to carry only that remainder, because the STA tool separately computes the switching energy of the external net from its own capacitance:

$$E_{int} = V_{DD} \int i_{VDD} dt - C_L V_{DD}^2 \quad (\text{rising edge}).$$

CharLib 2.1.0 does not measure internal power at all. The package contains no procedure writing `internal_power`, `rise_power` or `fall_power` — energy appears only as a unit definition — and `Characterizer.analyse_cell` schedules exactly: pin capacitance, then delay + leakage (combinational) or the constraint procedures (sequential). The shipped Liberty consequently has no internal-power groups; any power-aware flow consuming it sees switching and leakage power only.

lctime is the only tool of the three that writes power tables, inside the same transient as the delay measurement (`timing_combinatorial.py:263–274`):

```
gate_energy    = np.mean(gate_current * input_voltage) * dt
supply_energy  = np.mean(supply_current * vdd) * dt
switching_energy = gate_energy + supply_energy
```

Two defects are visible from first principles. First, the quadrature: $\text{mean}(f) \cdot (t_N - t_0)$ equals $\int f dt$ **only on a uniform time grid**; ngspice’s adaptive stepping clusters points around the edges, where $i v$ is largest, so the sample mean over-weights the transition region. The unbiased estimator on the same data is the trapezoidal sum

$\int f dt \approx \sum_k \frac{1}{2}(f_k + f_{k+1})(t_{k+1} - t_k)$, i.e. `np.trapz(f, time)` — available but unused. Second,

the accounting: no $C_L V_{DD}^2$ subtraction is performed, so the tabulated “internal” energy includes the full load-charging energy and grows linearly with C_L ; an STA tool then counts that energy a second time as switching power. Sequential cells get no power tables even from lctime (the supply current is retrieved in the FF harness but never integrated).

5.2 Leakage

Static power is a per-state quantity because subthreshold leakage depends exponentially on the bias of every device in a stack,

$$I_{leak} \propto e^{(V_{GS} - V_{th})/nV_T} (1 - e^{-V_{DS}/V_T}),$$

so a NAND2 leaks differently in each of its four input states (stack effect). CharLib does the canonical thing: one DC operating point per input combination, all 2^N of them, with

$$P_{leak}(\text{state}) = V_{DD} |I_{VDD}(\text{state})|$$

read from the supply branch (leakage_power.py:67, 86–91), emitting one conditioned leakage_power { when : ... } group per state. lctime's characterize_leakage_power is an empty stub, never called — no leakage data at all. For the sequential cells (which CharLib's combinational-branch leakage never reaches, and which have no unique DC operating point anyway), the project's seq_leakage.py measures a *transient average* instead: a reset pulse forces a known state, then $P = V_{DD} |\overline{i_{VDD}}|$ with the mean taken by .meas tran AVG over the settled window 40–60 ns — a documented single-state approximation.

6 Sequential constraints

6.1 The physics: pushout

As the data edge approaches the clock edge, a flip-flop's internal regeneration starts from an ever-smaller initial imbalance and the clock-to-Q delay diverges logarithmically (regeneration time constant τ_m of the cross-coupled pair):

$$t_{C2Q}(t_{su}) \approx t_{C2Q, \infty} + \tau_m \ln \frac{t_0}{t_{su} - t_{su}^*}.$$

A setup (hold) constraint is therefore not a hard edge but a chosen point on this curve, and every characterizer must pick a *criterion*: how much delay pushout $\Delta d = t_{C2Q} - t_{C2Q, \infty}$ is tolerable. Inverting the curve shows how the criterion maps to the constraint:

$$t_{su}(\Delta d) = t_{su}^* + t_0 e^{-\Delta d / \tau_m},$$

so a tighter (smaller) Δd yields a larger, more conservative setup time — but only logarithmically slowly, which is why differing criteria still produce comparable libraries.

6.2 CharLib: what is stock and what is local

Stock CharLib 2.1.0 is largely unimplemented here: the sequential *delay* procedure returns immediately (delay.py:14–15, a # TODO stub), as do recovery, removal, min-pulse-width and the metastability binary search. Its one working constraint procedure is the C2Q *contour* method (per the setup/hold-interdependence literature): reference C2Q at a relaxed corner, validity criterion $t_{C2Q} \leq 1.2 \times t_{C2Q, ref}$, exponential bracketing plus binary search for the

window edges, then an $N \times N$ sweep of the (setup, hold) rectangle and knee-point selection on the pass contour. On these HV cells the contour harness built transients that exhausted memory, so the shipped library uses the project's procedures (`seq_delay_procedure.py`, registered through CharLib's own registry):

- **clock→Q**: preload pulse, data switch mid-period, measured clock edge at the grid slew; ngspice `.meas` from the clock's 50 % crossing (second rise) to the output's 50 % crossing, 20–80 % output transition; tables indexed clock-slew $\times$ load.
- **setup/hold**: pass/fail **bisection** on a two-pulse harness, search range $-2 \dots +8$ ns, tolerance 10 ps. The pass criterion is deliberately two-part: the output must reach the target final state **and** its first half-rail crossing must occur within $1.5 \times$ the nominal C2Q ($\Delta d = 0.5 t_{C2Q, \infty}$, a *relative* degradation criterion) — final-state-only acceptance would silently ride through metastability. Latches are constrained against the closing enable edge; non-convergence counts as fail.

6.3 lctime

lctime's sequential path is genuinely implemented and more mathematically polished: it measures the pushout curve itself. `find_min_data_delay` doubles the setup/hold window until t_{C2Q} converges to $t_{C2Q, \infty}$ (abstol 1 ps); the constraint is then the root of

$$t_{C2Q}(t_{su}) - (t_{C2Q, \infty} + \Delta d) = 0,$$

with $\Delta d = 10$ ps (the default `max_pushout_time`), found with Brent's method (`optimize.brentq`, `xtol = 1` fs) after exponential bracketing — an *absolute* pushout criterion, in contrast to CharLib-local's relative one. It then re-solves setup with hold fixed at its unconditional minimum plus a 10 ps margin (and vice versa), which is what populates the constraint tables. Min-pulse-width uses the same root-finding on a two-pulse harness. Caveats visible in the source: the constraint search runs at **zero output load** (flagged # TODO in `flipflop.py`:196), `clk→Q` tables are taken at a fixed generous setup = hold = 1 ns rather than the measured constraints, and **latches are not supported at all** (`main_lctime.py` aborts) — where the shipped library needed `en→Q` and closing-edge constraints for its five latches.

For this library's C2Q of a few hundred ps, CharLib-local's $\Delta d = 0.5 t_{C2Q, \infty}$ ($\sim 100+$ ps) versus lctime's 10 ps sit far apart on the pushout curve, yet by the logarithmic inversion above the extracted constraints differ only by $\tau_m \ln(\Delta d_1 / \Delta d_2)$ — a few regeneration time constants, i.e. tens of ps. The bigger structural difference is coverage: only the local CharLib procedures produced latch data at realistic loads.

7 Numerics and simulator coupling

All three tools run the same engine — ngspice with the IHP OSDI-compiled PSP103 Verilog-A MOS models — through three different couplings, and two of the project's hardest-won findings are couplings, not physics:

	fo4.py	CharLib	lctime
coupling	ngspice -b batch	libngspice in-process (PySpice fork)	subprocess, ASCII wrdata
.meas support	native	fork-added; lost in subprocess mode	not used
integration method	trapezoidal (default)	trapezoidal, trtol = 1	trapezoidal (no options set)
timestep	20 ps fixed print step	slew/8 (delay), slew/10 (cap)	10 ps default, stop when breakpoints
stop condition	fixed 3 periods	autostop on .meas completion	breakpoint at 1 %/99 % V_{DD}
OSDI model loading	.spiceinit (batch honors it)	.spiceinit copied into run dir	.spiceinit copied by compare script
temperature	deck value	25 °C	hard-coded 25 °C (liberty header ignored)

The trapezoidal rule's local truncation error per step, $\epsilon \approx \frac{h^3}{12} |d^3v/dt^3|$, is what ngspice's timestep control bounds; trtol scales the tolerance it is compared against, so CharLib's trtol = 1 (versus the default 7) forces $\approx 7^{1/3} \approx 1.9 \times$ finer steps through the transitions — a deliberate accuracy/runtime trade in the delay decks.

The two coupling landmines, both documented in run_charlib.sh and the config generator because both produce *silently* wrong libraries: (1) with the ngspice-subprocess backend PySpice parses only the rawfile and never reads measurement results back, so every delay table comes back empty while leakage and capacitance still populate — the file looks plausible until a timing tool finds no arcs; shared mode is mandatory. (2) ngspice reads .spiceinit in batch and interactive mode but *not* in server mode (ngspice -s), so PySpice-driven runs never load the OSDI PSP103 models without the project's pre_osdi shim. lctime's file-based batch backend is immune to both by construction — its weakness is instead that it sets no simulator options at all and inherits whatever .spiceinit the working directory supplies, and that it stamps .option TEMP=25 regardless of the library header (the nom_temperature it parses is never propagated; characterizing a non-25 °C corner silently produces 25 °C data).

8 Can lctime use the charge-integration method?

As shipped in 0.0.26 — no. The verdict rests on three code facts:

1. The input-capacitance procedure is irreducibly the constant-current slope method: it drives the pin from a current source, records *only* the pin voltage (output_currents=[], input_capacitance.py:253–257), and computes $C = I \Delta t / \Delta V$ from two threshold crossings (lines 273–297). There is no charge in sight.

2. No code path in the package integrates a current: a search over the installed tree finds no `trapz/cumsum` and no `ngspice integ` in any generated control block.
3. The 10 μA source magnitude is a hard-coded constant (`util.py:162`), not reachable from the CLI — so even the existing method cannot be tuned, let alone replaced, by configuration.

But the port is small, because the data path already exists. For its power tables `lctime` already requests branch currents through named voltage sources and gets them back aligned with the time vector — both backends implement it (`wrdata ... i(vpin)` file-based; `print i(vpin)` interactive) — and the combinational routine already retrieves the *gate* current `-currents["V<pin>"]` for its energy term. Charge-integration capacitance is the same measurement pointed at a different source. The patch, confined to `characterize_input_capacitances()`:

1. **Drive voltage, not current:** declare the active pin as a voltage input and replace the `I<pin>` PWL current source with the `StepWave` ramp the delay code already uses (using the Liberty-correct $T = s/0.6$ stretch while at it).
2. **Ask for the current:** `output_currents=["V<pin>"]` instead of the empty list.
3. **Integrate:** replace the secant-slope block with

$$Q = \int_{\text{edge}} i_{\text{pin}}(t) dt \approx \text{np.trapz}(-i_{\text{Vpin}}, t), \quad C = \frac{|Q|}{V_{DD}},$$

using the trapezoidal sum (not the rectangle rule the power code uses) since the time grid is adaptive.

Nothing else in the harness changes — grids, state enumeration, reduction, and Liberty emission all operate on the returned scalar. On the evidence of this library, the payoff is the difference between +52 % and +9.7 % on `inv_1`'s pin capacitance, i.e. between a wire-load estimate that is wrong by half and one within the Miller-charge ambiguity that any single-number capacitance carries. (`lctime` is AGPL-3.0-or-later; a modified copy used in-house carries no distribution obligation, and an upstream contribution would resolve the license question entirely while fixing the rise/fall label swap noted in §3.)

9 Provenance of the shipped Liberty: no `lctime` data

A natural question given three tools: which of them actually populated `lib/sg13g2_stdcell_hv_typ_3p30V_25C.lib`? The answer is that every number in the shipped file traces to the `CharLib` flow, and **none to `lctime`**:

Liberty content	produced by
combinational <code>cell_rise/fall</code> , <code>rise/fall_transition</code> tables	<code>CharLib 2.1.0</code> (<code>run_charlib.sh</code>), post-processed by <code>fix_lib.py</code>
pin capacitance / <code>rise_capacitance</code> / <code>fall_capacitance</code>	<code>CharLib</code> charge_integration procedure

Liberty content	produced by
per-state leakage_power groups (combinational)	CharLib all-states DC enumeration
clk→Q and en→Q arcs, setup/hold constraint tables	local seq_delay_procedure.py inside CharLib's registry, merged by merge_lib.py, fixed by fix_lib_seq.py
sequential and tie-cell cell_leakage_power	local seq_leakage.py / tie_leakage.py (transient average)
slew and load index grids	thin-oxide grids × the fo4.py ratios (2.66 delay, 2.20 capacitance)
functions, pin directions, ff/latch groups	lifted from the thin-oxide Liberty (logic unchanged by the transform)

lctime's role was strictly read-only verification: lctime_compare.py re-characterizes eight combinational cells on the same grids and models and compares its output *against* the shipped tables — the 2.9 % median delay agreement, the −19 % slew-convention band and the +52 % pin-capacitance finding of the preceding sections are all products of that comparison, and none of its data flows back. Given what the comparison revealed about the slope-method capacitance (§3) and the slew convention (§4), that separation is the correct engineering outcome, not an accident of history. A text search of the shipped file confirms it carries no lctime or LibreCell traces.

10 Cells outside every characterizer's model

Nine of the 84 cells are drawn but were not characterized by either tool, and in every case the obstacle is the characterizer's data model rather than anything about the cell. The three mechanisms are worth separating, because only one of them announces itself.

No high-impedance state (the 6 tri-states). CharLib's cell schema (charlib/config/syntax.py:30-120) has keys for inputs, outputs, functions, clock, enable, set, reset, state — and nothing for three_state or a Hi-Z output value. A search of the entire installed package for high.?impedance|high-?z|hiz|three.?state returns exactly one hit, a comment on Port.Role.ENABLE. So a tri-state cell cannot be expressed. What makes this the dangerous case is that gen_charlib_config.py does *not* skip these cells: Z has a function (A or !(A)), TE_B is collected as an ordinary input, and the cell is configured as a two-input combinational gate whose output ignores one input. CharLib then emits an empty result, omit_on_failure swallows it, and **the run log never mentions the cell**. This is the same silent-failure class that once hid all 14 sequential cells.

No statetable (the 2 clock gates). lgcp_1 and slgcp_1 hold state in a statetable with an internal pin, not in an ff/latch group. CharLib has no input form for that, so gen_charlib_config.py:136-142 skips them explicitly, with a stated reason — the honest failure mode.

No output pin (sighold). The bus holder's only signal pin is an inout, so the configurator's "no output pin" branch drops it, together with the genuinely arc-less fill/decap/antenna cells.

lctime does not rescue any of them. It *is* tri-state aware — it reads `three_state` into `CombinationalOutput(high_impedance=...)` and writes the attribute back out — but its characterization deliberately pins the enable to its active value (`constant_input_pins`, "Skip measuring the tri-state enable input") and characterizes only the data arc. The enable arcs, which are most of what a tri-state Liberty group contains, are exactly what it omits. For clock gates it has nothing: no `clock_gating_integrated_cell`, no `statetable`, and its sequential recogniser bails on tri-state outputs outright.

The answer is the one the tie cells and the sequential constraints already use: measure directly against the shipped netlist with ngspice, and emit the Liberty groups from the measurements.

10.1 A fourth answer to the capacitance question: the bus holder

Section 3 compared three ways of averaging $C_{in}(v)$. `sighold` adds a case where one of them is not merely different but **invalid**.

Charge integration — the production method for gate pins, and the defensible one there because it measures the charge a driver actually delivers — assumes the integrated current is the charge that ends up on the pin capacitance. On a bus holder it is not: the keeper actively fights the driver until the internal inverter flips, and that crowbar charge lands in the same integral. It is therefore not a property of the cell. Measured on the shipped netlist, integrating the driven edge:

driving edge	0.2 ns	0.5 ns	1 ns	2 ns	5 ns
rise charge	30.5 fC	38.1 fC	46.3 fC	58.4 fC	86.4 fC

A monotone 2.8× spread across a plausible range of driver strengths. Any single point reported as capacitance would silently encode an arbitrary assumption about whatever drives the net.

The rail-biased AC method of section 3.2 has no such coupling, because it never asks the keeper to lose: at a fixed bias the keeper is in a defined state and the small signal sees only the physical capacitance. Four bias points near the rails give 0.00377 / 0.00378 / 0.00353 / 0.00361 pF — a 7 % spread, and consistent with the ~2.1 fF of gate area plus junction contributions. Mid-rail is excluded on purpose: the cross-coupled pair has gain there and the small-signal result is meaningless.

This also explains the thin-oxide library's `sighold`, which reports 0.0268 pF rising against 0.0096 pF falling — a 2.8× asymmetry with no structural cause, since the physical capacitance of a node has no direction. It is a fight-charge artifact of whatever transient method produced it. The thick-oxide cell therefore ships equal rise and fall values, with the keeper fight-back charge documented separately as the driver-dependent quantity it is.

The leakage numbers land where the rest of the library's do: 0.0118 nW holding high and 0.0224 nW holding low, roughly 10^4 below the thin-oxide cell, the same ratio the tie cells show and for the same reason — 0.45–0.70 μm channels under a thick oxide.

10.2 Tri-states: three arcs, and one that cannot be measured as specified

A tri-state cell needs three arc classes rather than one: the combinational data arc — which either tool could produce, and which is measured here exactly as for any other gate — plus `three_state_enable` and `three_state_disable`, which are distinguished from ordinary delays by measuring *into* and *out of* a floating state. `ebufn` carries four tables per enable arc; `einvn` uses the `_rise` variants and carries two. All six cells are characterized by `work/char_tristate.py`; the harness is calibrated by re-measuring a shipped cell's pin capacitance (`sg13g2_hv_buf_4` pin A) through the same code path and requiring agreement with the shipped library — it lands at **0.46 %**.

The **enable** arc is straightforward once the floating node is defined: the output starts in Hi-Z held by the 1 G Ω mid-rail keeper the functional suite already uses for its 12 Hi-Z checks, `TE_B` switches, and the measurement is the usual 50 %-to-50 % delay and 20/80 transition into the specified load.

The **disable** arc is where the specification breaks down. The obvious recipe — wait for the output to leave its driven level through the slew threshold — is not physically measurable on a capacitively loaded Hi-Z node. Once the driver releases, nothing but the keeper moves the node, so the answer is $0.2 RC \approx 0.5$ ms and is *exactly proportional to the load*. No keeper value repairs this: to make the decay comparable with the turn-off time it would have to approach the cell's on-resistance (~ 1 k Ω), at which point the keeper fights the enabled driver and the cell can no longer hold its own output.

The thin-oxide library cannot have measured it that way either, and its own tables say so: its disable entries are *exactly* load-independent, the seven values along the load axis differing only by a 1 fs increment — a monotonicity epsilon, not a measurement.

What is measured instead is the quantity that actually goes away: the drive current. Z is held by an ideal source at the slew threshold it would cross on leaving the driven level (0.66 V driving low, 2.64 V driving high), the on-state drive current I_{on} is recorded (0.56 / 0.60 mA for `ebufn_4`), and the arc is the time from `TE_B` crossing 50 % to $|I_Z|$ falling through $I_{on}/2$ — the same 50 % convention used everywhere else in this library, applied to current rather than voltage. This is load-independent by construction, so the seven rows are identical with **no synthetic epsilon**, and the transition tables repeat the delay because nothing traverses a slew after release (as the thin-oxide library also does).

10.2.1 What the ratios say

cell	comb. c_r/c_f	enable c_r/c_f	disable c_r/c_f
<code>ebufn_2</code>	2.22 / 3.01	2.02 / 1.92	1.48 / 2.10

cell	comb. c_r/c_f	enable c_r/c_f	disable c_r/c_f
ebufn_4	2.21 / 2.93	2.01 / 1.87	1.46 / 2.03
ebufn_8	2.20 / 2.93	2.00 / 1.80	1.52 / 1.87
einvn_2	2.11 / 2.74	2.00 / –	1.47 / –
einvn_4	2.11 / 2.74	2.00 / –	1.49 / –
einvn_8	2.09 / 2.73	1.98 / –	1.53 / –

Median over load points 2–7 against the thin-oxide library; the expected band is the library’s 2.66× delay ratio $\pm 20\%$. Load point 1 is excluded because the thin-oxide tri-states offset their load axis by the Z pin capacitance (0.0098 pF rather than 0.001), so at that point the HV cell carries a 10× lighter load than its reference.

The **combinational** arcs are in band. The **enable** delays sit at 1.8–2.0, just under it — but the denominator is the problem, not the measurement: the thin-oxide enable tables are *clamped* at their fast end (ebufn_4 repeats 0.074755, 0.074756, 0.074757, 0.074758 — a 1 fs floor rather than four measurements). The enable *transitions*, which are the same physical quantity as the combinational transitions that do agree, come out at 2.25–2.28, in band.

The **disable** ratios (1.46–1.53) are out of band **by construction**, because the criterion differs from whatever produced the thin-oxide numbers. They are reported rather than tuned. Their tight clustering across two footprints and three drive strengths is what a consistent criterion offset looks like; noise would not cluster.

10.2.2 Two infrastructure defects found on the way

Worth recording because both would have produced plausible numbers:

- The mid-rail keeper used to define the floating node was left attached during **leakage** measurement, where it sourced its own 1.65 nA into the supply reading — a fictitious 5.48 nW on ebufn_2 in the A & !TE_B state, **150× the true value**. The keeper now exists only where the node actually floats. A leakage number that is 150× wrong is still a perfectly plausible-looking leakage number.
- ngspice must run with OMP_NUM_THREADS=1 here. Its spinning OpenMP barriers burned 130+ CPU-seconds without completing on a contended machine, against ~4 CPU-seconds to completion single-threaded.

10.3 Clock gates

The clock gates need a CLK→GCLK propagation arc (which the thin-oxide library writes with *no timing_type* at all, i.e. combinational), setup_rising/hold_rising on the enable pins against the clock — the form the local bisection procedure of section 7.2 already emits — and a min_pulse_width constraint on the clock pin, obtained by bisecting the clock pulse until the gated output fails to produce a full-swing pulse. CharLib has a config slot for min_pulse_width_constraint_procedure but never calls it; its dispatch is a # TODO.

11 Drive limits: the attribute whose absence fails twice

CharLib emits no `max_capacitance` and no `max_transition`, on any pin, in any cell — and no library-level `default_max_*` either. The shipped library inherited that gap for three revisions, and it fails in two very different ways.

Silently. OpenSTA answers “no limit” for every pin, so a flow’s max-capacitance and max-transition checks report zero violations. They are not passing; there is nothing for them to compare against. A green signoff report is the worst possible presentation of an absent constraint.

Loudly, and much later. OpenROAD’s TritonCTS consults the same data when it sizes the clock tree, gets an empty buffer selection back, and dereferences it: `clock_tree_synthesis` terminates with SIGSEGV inside `cts::TritonCTS::getBufferFanoutLimit`. The failure surfaces in a different tool, several flow stages after the actual defect, with a stack trace that says nothing about liberty. Reduced to a two-flip-flop test case it reproduces deterministically, and it disappears the moment the limits are present — reported upstream as OpenROAD issue #11165, on the argument that a liberty without limits deserves a diagnosable error rather than a signal 11.

The limits are derived from the characterization rather than borrowed: a pin may not be asked to drive more load than its tables cover, nor to accept a slower edge than was characterized. So per output pin `max_capacitance` is the top of that pin’s load axis, per pin `max_transition` the top of the slew axis it was characterized against, and the library defaults follow the thin-oxide convention of weakest-drive / slowest-edge. These bound the characterized range; real slew targets belong in the flow constraints, not here.

12 Conclusions

1. **All three input-capacitance methods are averages of the same $C_{in}(v)$ over different windows**, and their results order exactly as the windows predict: rail end-points 5.87 fF < full-swing charge mean 6.44 fF < Miller-peak window mean 8.92 fF. Charge integration is the defensible production choice — it measures the charge a driver actually delivers — with the rail-biased AC value as the clean lower anchor.
2. **CharLib is Liberty-correct on the slew convention; lctime is not.** The $T = s/0.6$ ramp stretch is the single largest systematic between the two characterizers (–19 % at multi-ns slews, mechanism confirmed by re-interpolation at 0.6 s to 0.25 %); where the convention doesn’t bite, the tools agree to a median 2.9 % on delays and 0.0 % on transitions.
3. **Neither tool covers the full Liberty power model.** CharLib 2.1.0 has no internal-power procedure at all; lctime writes power tables but with a biased quadrature and without subtracting the $C_L V_{DD}^2$ load term, so they double-count against STA switching power. Leakage is CharLib-only (all-states DC enumeration; the project adds transient single-state leakage for sequential cells).

4. **Sequential characterization shipped only because of local code:** upstream CharLib's sequential delay is a stub and its contour method was unusable here; the local bisection (relative 1.5× C2Q criterion) and lctime's Brent-on-pushout (absolute 10 ps) are both principled points on the same metastability curve, differing mainly in coverage — lctime cannot do latches and constrains at zero load.
5. **lctime can adopt CharLib's charge-integration method with a three-part, single-function patch** — the branch-current infrastructure it needs is already exercised by its own power measurement. Until then, its pin capacitances should not be used for anything load-sensitive.
6. **Neither characterizer models tri-states, clock gates or bus holders**, and the tri-state case fails *silently* — the cell is configured, no simulation runs, an empty result is swallowed, and the log says nothing. Any characterization flow needs a coverage check that compares the emitted cell list against the intended one; a clean run log does not mean the run was complete. These nine cells are measured directly against the shipped netlists instead.
7. **A method can be correct for one cell class and invalid for another.** Charge integration is the right production choice for gate pins and the wrong one for a bus holder, where the integral is dominated by keeper fight-back and varies 2.8× with the driver's edge rate. The bias-swept AC method, which is only the *anchor* for gate pins, is the *production* method there. The thin-oxide library's unexplained 2.8× rise/fall asymmetry on the same cell is the same artifact, visible in shipped data.
8. **max_capacitance and max_transition are not optional.** Absent, they make STA limit checks pass vacuously and crash OpenROAD's CTS in a different tool several stages downstream. Deriving them from the characterized table axes costs nothing and is self-consistent by construction: a cell is never asked to operate outside the range it was measured over.

13 Appendix: provenance

item	identity
reference deck	work/fo4.py, 96 lines, ngspice batch
CharLib	2.1.0, git stineje/CharLib commit 6859faf, venv /foss/tools/charlib, install hash-verified unmodified
PySpice	1.6 fork, git infinitymdm/PySpice commit da81c4d (adds .meas readback)
lctime	0.0.26, /usr/local/lib/python3.12/dist-packages/lctime
models	IHP SG13G2 PSP103 Verilog-A via OSDI, cornerM0Shv.lib mos_tt, 3.3 V, 25 °C
shipped library	

item	identity
	lib/ sg13g2_stdcell_hv_typ_3p30V_25C.lib, thresholds 20/80/50, 7×7 NLDM grids
cross-check	work/lctime_compare.py: 8 cells, 3 132 aligned points
local CharLib adaptations	charlib_patched.py (case-insensitive branch lookup, procedure registration), seq_delay_procedure.py (clk→Q, setup/hold bisection), seq_leakage.py, gen_charlib_config.py (grids ×2.66/×2.20, charge integration selected), fix_lib.py / fix_lib_seq.py (header and sequential emission repairs)
direct measurement, outside both characterizers	tie_leakage.py (tie cells), char_sighold.py (bus holder: settled-tail leakage, bias-swept AC capacitance, fight-charge sweep), char_tristate.py (data + enable/disable arcs), char_clockgate.py (CLK→GCLK, setup/hold, min pulse width)
Liberty post-processing	finalize_lib.py (drive limits from the table axes, physical-cell stubs with measured leakage, corner-suffixed library name)
Liberty gate	verify_lib.py: structure, cross-view, areas vs layout, load-axis monotonicity, C _{in} vs reference, sequential arcs, drive limits, and complete- construct checks for the tri-state / clock-gate / bus-hold classes

Key source locations cited: CharLib procedures/combinational/delay.py (stimulus 75–83, measurements 101–112, worst-case 188–190), pin_capacitance/charge_integration.py (integration 100–104, formula 124–135), combinational/leakage_power.py (67, 86–91); lctime characterization/timing_combinatorial.py (StepWave 186–191, energy 263–274), characterization/input_capacitance.py (current source 235–243, slope 273–297, label swap 318–321), characterization/timing_sequential.py (pushout root-finding 1018–1228), ngspice_subprocess.py (deck 122–187, temperature 83).